

JD Bank — HR Review & Decision Matrix

For: SFU Human Resources.

Purpose: everything HR needs to sign off on how the JD Bank evaluates a job description — what the system is, how it works, what we measured against your real job descriptions, and the specific settings that are yours to rule on. This document is self-contained; **you do not need to read any code**, and every decision below is a setting we change in minutes and re-measure.

Status: *awaiting HR review — nothing in this system has been ratified by SFU.* Until you rule, the system's behaviour is provisional and unchanged.

Evidence base: every number here is measured against the **entire SFU job-description archive — 14,565 files**. It is not a sample.

We are asking you for two things:

1. **Rule on eight settings** (Part 4) that decide whether a job description can be approved.
2. **Do a short hands-on pilot** (Part 5) — walk a handful of real job descriptions through the tool the way a reviewer would, and tell us where it got the call wrong.

Together this is how the JD Bank moves from *built* to *trusted*. Budget about 60–90 minutes.

1 · What the system is

The JD Bank is a tool that helps SFU **standardise, quality-check, and author** APSA / APEX / Polytechnic job descriptions against SFU's own JD standards. It does three jobs:

1. **Reads a job description into its sections** — Position Summary, Duties, Qualifications, and so on.
2. **Checks it against SFU's JD standards** — a rulebook of gates, word lists and thresholds. The parts SFU has published come from the **SFU JD Toolkit** and the official **APSA/APEX/POLY job description template**; the parts the Toolkit is silent on are settings we chose (and this document asks you to ratify the ones that matter).
3. **Produces a result** — a **score out of 100**, a **grade A–F**, a **list of specific, located issues**, and a decision on **whether an HR reviewer is permitted to approve the job description**.

It supports two ways of producing a job description, both ending at the same human-approval gate:

- **Harmonising the archive** — where several near-duplicate job descriptions describe the same role, the tool merges them into one clean **draft** canonical job description and shows exactly what each source contributed and what was dropped.
- **Authoring a new one** — a hiring manager or recruiter answers a guided set of questions (or starts from an existing role) and sees the JD scored live against SFU standards as they write.

The one rule that never bends: the system never approves anything — a human always does. The tool can *recommend*, and it can *block the approve button*, but publishing a job description is always a deliberate human act, recorded in an audit log.

Where the numbers in the rulebook come from

To evaluate a job description the system uses **194 settings** — numbers, word lists and rules.

Source	Count	What it means
SFU's own JD Toolkit	19	Published SFU standards (e.g. the 100–150-word Position Summary range).
Our choices	175	Settings we had to pick because the Toolkit does not specify them.

Most of the 175 govern the machinery (how duplicates are detected, how roles are clustered) and change nothing about whether a JD is acceptable. **A small core of them silently decide which of your job descriptions pass — and those are what this document puts in front of you.** The remaining ~185 are recorded internally and, the archive shows, change almost nothing; they can be ratified in bulk or reviewed at your leisure.


How to read a result

Element	What it means
Score / 100	Overall quality. Higher is better.
Grade A–F	A banded view of the score.
Issues	Specific, located findings — e.g. "Summary is 172 words; the range is 100–150."

Element	What it means
Gates	The subset of issues that can block approval . Most gates can be overridden by a reviewer who writes down a reason; a few cannot (see Decision 5).

2 · How the process works

A job description moves through the bank in one direction, and a human stands at the only exit:



Five guarantees are built into this flow. They are **not** up for debate — they are the safeguards the decisions in Part 4 sit inside:

- **Nothing auto-publishes.** Every canonical job description is a *draft* until a reviewer approves it.
- **Overrides require a written reason.** A reviewer can waive most gates, but the reason is recorded.
- **The audit log is append-only.** Every approve / reject / edit / override, and every sign-in, is traceable to a person and cannot be altered after the fact.
- **Every canonical job description traces back to its sources** — which archive JDs fed it, and why content was kept or dropped.
- **All processing runs on SFU-controlled infrastructure.** No job-description text is sent to any outside vendor.

3 · What we measured

The point of the baseline run was to put the proposed approval bar on trial against SFU's **real** job descriptions — and find out whether it is reasonable or too harsh.

Read the right population

The archive spans decades and **two different job-description forms**. Three figures matter, and only one is a fair test of the bar:

Population	Size	Approvable	What it actually measures
Whole archive	14,522 scored	~5%	Not meaningful — do not quote. Mixes 1990s JDs and a second form (the CUPE Weighted Job Questionnaire) that has no SFU quality bar at all.
JDs dated 2024+	1,034	61%	A <i>date</i> band — still partly measuring a rollout, not quality.
Current practice	874	78.6%	The fair trial — the job descriptions SFU writes today, under today's template and conventions.

The CUPE form (4,300 files, ~30% of the archive) is the Weighted Job Questionnaire — a *different instrument* with no SFU quality bar defined. The tool reads it, but does **not** score it against the APSA/APEX/Polytechnic rules; doing so would be a category error. It is excluded from every approval figure above. Decision 8 asks you to confirm this scope.

The headline: the bar is sound, and barely binding

On the 874 job descriptions SFU writes today:

Would be approvable	687 — 78.6%
Median quality score	79 / 100
Grade spread	81 A · 551 B · 240 C · 2 D · 0 F
Clear the score floor of 60	99.8%
Rejected because their score was too low	2

Nearly eight in ten current job descriptions would pass straight through. The rest need **edits, not rewrites. No current job description scores an F.**

What actually blocks a job description

This is the single most important finding for HR, because it is **not** what most people assume. Of the 187 current job descriptions that cannot yet be approved, here is what stops them:

What blocks it	JDs	
Position Summary is outside 100–150 words	134	← the real gatekeeper
Missing "or an equivalent combination of education and experience"	42	
Missing part of the territorial / equity footer	10	
Summary describes working conditions, not the role	9	
Missing a mandatory section	7	
An issue's severity is too high	7	
The quality score itself being too low	2	
Grade too low	2	
Qualifications listed out of order	1	

The thing doing the real gatekeeping is a word count, not the quality score. The quality score rejects two job descriptions out of 874 — it is almost inert. That is not a flaw, but HR should ratify the bar *knowing* it, not assuming the score is what protects quality. (Decisions 1 and 3 both turn on this.)

4 · The decisions

First — three things we already corrected, so you are reading honest numbers

The first version of this review exposed **three defects in our own rules**. We fixed them and re-measured **before** bringing you this document — handing you figures we already knew were distorted, collecting your sign-off, and *then* fixing them would have made your ratification meaningless. **No action is needed on these three; they are here so you know why the numbers look the way they do.**

Was	What it was	Now
A qualifications rule scanned the whole document	It blocked 104 JDs for "wish-list" phrasing found anywhere — even a duty that said "responsibilities <i>may include...</i> "	Scoped correctly to the Qualifications section: 104 → 0 wrong blocks; 59 JDs became approvable. This is the entire improvement from 71.9% to 78.6%.
A "how and why" rule that could never not fire	It penalised every duty of every JD for missing detail the reader never extracts — an invisible constant subtracted from every score	Retired until the reader can extract that field. Median score rose 77 → 79 ; no approval changed (it only affected scores, not the pass/fail line).
Our "era" model conflated two rollouts	It judged correctly-written 2019 JDs by a footer rule only 2023+ JDs could satisfy — a sevenfold distortion, all date and no quality	Added a fourth "current" band from 2024. The 78.6% figure is measured on the corrected model.

Now — the eight settings that need your ruling

Each is a **setting**, not a code change. Three kinds of ask appear:

- **Ratify** — we recommend keeping it; we need you to *own* it knowingly.
- **Choose** — a genuine policy fork; pick an option.
- **Review** — needs an experienced JD reviewer's eye, not an engineer's.

#	Decision	Type	Our recommendation	Blocks today	Your ruling
1	Position Summary must be 100–150 words	 Ratify	Ratify — it's SFU's own number, but know it is the #1 determinant of approval	134	☐
2	Too- short summaries are not blocked (only too-long)	 Choose	Keep the asymmetry, deliberately	0	☐
3	Minimum bar: score ≥ 60 · grade ≥ C · severity ≤ high	 Ratify	Ratify all three (they reject 2 of 874); 70 is available if you want a real bar	2	☐
4	Territorial + Employment-Equity footer is mandatory	 Choose	Auto-insert the footer at compose time instead of penalising omission	10	☐
5	The " placeholder " gate blocks with no appeal	 Choose	Make it waivable with a written reason	0	☐
6	The wish-list-language word list ("may include", "an asset")	 Review	A JD reviewer confirms the list is complete	0	☐
7	Should the wish-list gate be overridable ?	 Choose	Confirm it stays waivable (decide on purpose)	0	☐
8	Scope: the tool scores/authors only APSA/APEX/Poly, not CUPE (~30% of the archive)	 Choose	Confirm APSA/APEX/Poly-only for now, or commission a CUPE bar (a separate project)	n/a	☐

Details for each follow.

Decision 1 — Ratify the 100–150-word summary range

The question: SFU's Toolkit says a Position Summary should be 100–150 words. We enforce that as a gate. Do you ratify it?

What you need to know first: this single rule is **the largest determinant of whether a job description is approvable** — it blocks 134 of the 187 current JDs that fail. It is not a formatting nicety; it is the operative bar. The saving grace is that **100–150 is SFU's own published number**.

Recommendation: ratify as-is, with eyes open. If you'd rather widen it (e.g. 100–200), we change one value and re-measure — expect the 134 blocks to fall sharply.

Decision 2 — Keep the short-summary asymmetry

The question: today a summary that is **too long** is blocked, but one that is **too short** only loses points — it is not blocked. The Toolkit's range is two-sided; our enforcement is one-sided. Make it symmetric?

What you need to know: **340 of 874** current job descriptions have summaries that are *too short*. If we blocked those too "for consistency," the under-run would instantly become **the single biggest blocker in the system**, and approval would fall from ~79% to roughly **40%**.

Recommendation: keep it as-is, deliberately. We flag it only so the asymmetry is a decision you made, not an accident. If you do want symmetry, we will show you the ~40% figure first and have you ratify *that number*, not the principle.

Decision 3 — Ratify the minimum quality bar

The question: a job description cannot be approved if its score is below **60**, its grade below **C**, or it carries an issue of severity above **high**. These three thresholds are ours — SFU publishes no such numbers. Ratify them?

What the archive says: they are nearly inert on current practice. 99.8% of today's JDs clear the score floor; **it rejects two**. The bar survived its trial against real data.

Recommendation: ratify all three. Two honest caveats:

- The floor of 60 is defensible because it is **nearly inert**, not because 60 is a magic number.
- If SFU wants a **more demanding** bar later, the data supports it: the median current JD scores **79**, and 81 score an A. **A floor of 70 would be a real bar** rather than a formality. That is a policy choice, cheap to make — one value and a re-measure. We are not recommending it, only

telling
you the option is there.

Decision 4 — How to handle the territorial / equity footer

The finding: the mandated territorial acknowledgement and Employment-Equity statement are required by the rulebook. **94% of the archive lacks them** — but this is **not** a quality problem. SFU only began putting the acknowledgement into job descriptions around 2023:

Year	JDs carrying the acknowledgement
2019	0.2%
2023	11%
2024	63%
2025	85%
2026	89%

So the rule is really detecting **a document's age**, not its quality. Staff did nothing wrong; the paragraph simply did not exist yet. The gate *can* be waived — but on the legacy archive that means a reviewer writing an individual justification roughly **13,000 times**.

Choose one:

- **(a)** Accept as-is. Old JDs stay blocked until someone adds the paragraph. *(Honest, but makes the historical archive practically unusable without ~13,000 waivers.)*
- **(b) [recommended]** Treat the footer as **boilerplate the system inserts automatically** when a JD is composed, rather than something the author is penalised for omitting. It is identical on every JD. *(If we generate the wording, HR must give us the current official text to insert.)*
- **(c)** Apply the rule only to JDs written after a date you choose.

Recommendation: **(b)**.

A related refinement, if you want it: we currently call a JD "current practice" by its *date* (2024+). The truer signal is whether it actually *carries the footer*. The two nearly agree, and either is defensible — this is a definitional call about what "current practice" means at SFU, and it is yours. Not required to proceed.

Decision 5 — Make the "placeholder" gate appealable

The finding: most gates can be waived by a reviewer with a written reason. **Two cannot** — a job description that trips them is permanently un-approvable, with no human override. One of those two is demonstrably wrong.

The placeholder gate treats the phrases "**action verb**," "**how and why**," and "**what by**" as evidence that template boilerplate was left in the document. But a job description that legitimately *discusses* action verbs — as a **Communications or HR training role might** — is then permanently un-approvable, and no reviewer who can see the JD is fine has any way to say so.

Good news: it blocks **zero** current job descriptions. It is a legacy-archive issue (29% of old files), not a threat to what SFU writes today.

Recommendation: remove the placeholder gate from the no-appeal list. Keep the rule — it does catch real leftover template text — but let a reviewer waive it with a written reason. Removing human discretion is only defensible when a rule is *never* wrong, and this one provably is. (*The other no-appeal gate — "a mandatory section is missing" — is defensible; we recommend leaving it.*)

Decision 6 — Review the wish-list-language word list

The question: the rulebook blocks "wish-list" phrasing in Qualifications — words like "*may include*" or "*an asset*" that quietly turn a minimum requirement into a nice-to-have. Correctly scoped to the Qualifications section, that list now matches **only 10 files in the entire archive**.

That leaves one question **an engineer cannot answer**: is the list a guard-rail authors rarely trip — or is it **missing the phrases SFU's authors actually write?**

Recommendation: a few minutes of an experienced JD reviewer's time to look at the list and tell us whether it captures the wish-list language you see in practice. If it's missing phrases, we add them.

Decision 7 — Confirm the wish-list gate's override status

The question: the wish-list gate is currently **overridable** (a reviewer can waive it). We would like you to confirm that on purpose rather than inherit it.

Recommendation: keep it overridable unless SFU considers wish-list phrasing a hard stop, in which case we make it non-appealable. A small policy call, bundled here so it isn't decided by default.

Decision 8 — Confirm the scope: APSA/APEX/Poly only, not CUPE

The finding: everything above is measured on job descriptions written on SFU's **APSA/APEX/Poly** template. **CUPE** roles use a different form — the **Weighted Job Questionnaire (WJQ)** — and they are

4,300 of your 14,565 files (~30%). The tool *reads* them, but does **not** score them and will **not** author them.

This is deliberate, not an oversight:

- The whole quality bar in this document was built from SFU's APSA/APEX/Poly standard. **There is no CUPE quality standard encoded** — the Toolkit does not define one for the WJQ form.
- Scoring a CUPE job description against these rules would be a **category error** — a number that looks like a quality grade but really measures "how much does this CUPE role resemble an APSA/APEX/Poly one." That is worse than no number.

Choose one:

- **(a) [recommended, for now]** Confirm the Bank is **APSA/APEX/Poly-only**. CUPE JDs remain readable/searchable but are neither scored nor authored, and this document's numbers cover ~70% of the archive **by design**, honestly labelled.
- **(b)** Commission a **CUPE quality bar** — a WJQ ruleset with its own thresholds and sign-off. That is a **separate project**: it needs SFU to define what "good" means for a WJQ role before we can build a checker for it. Once that exists, CUPE switches on with one setting.

Either way, decide it **on purpose** rather than inherit "the Bank ignores a third of SFU roles" as an accident of which template we started with.

5 · Walk the system (the hands-on pilot)

Your JD Bank operator will send you the link and confirm your sign-in (your SFU account) and that you have the **reviewer** role. Once in, use the navigation bar at the top. The goal is to exercise a real reviewer's judgement on a handful of drafts and flag anything the tool gets wrong.

1. **Open**  **Review queue.** You'll see the drafts awaiting review, most-blocked first, each with its score and how many gates block it. Pick one to open.
2. **Read the draft page.** At the top: the **score, grade**, and whether it's *approvable* or how many gates block it. Below: the **blocking gates** (each says what tripped and whether it can be waived), the **draft text**, and a **"removed content" list** (what the tool folded out when it merged duplicate JDs, and why).
3. **Click "Changes since last approved version →."** If this role was approved before, you'll see a side-by-side of what changed. (First-time drafts say "no prior approved version" — expected.)
4. **Make a decision** on a few drafts, trying each action at least once:
 - **Approve** — if a gate is *waivable*, you'll be asked to type a reason before it lets you. A draft with a non-waivable blocking gate cannot be approved (by design).
 - **Reject** — requires a written reason.
 - **Edit** — change any field (title, summary, duties, qualifications, the footer flags...) and Save. This is a structured, field-by-field editor — no raw data. Saving creates a **new draft version** for review; it does **not** publish.
5. **(Optional) Try**  **Builder.** Author a brand-new JD from the guided questions (or search an existing one and clone it), press **Check compliance**, and watch the live panel score it as you go. Submit it and it lands in the same review queue.
6. **Keep a short list as you go:** any JD where you *disagreed* with the tool — it blocked something fine, or passed something you'd reject. **That list is the most valuable thing you produce.**

( *Dashboards* and  *Guide* are there for context — the dashboards summarise the archive; the guide

is the full operator manual. Neither is required for this review.)

6 · Recording your decisions, and what happens next

What we need back:

- **Your eight rulings** (Part 4) — keep or change, in a sentence each.

- **Your pilot notes** (Part 5) — which drafts you approved / rejected, and every case where you disagreed with the tool's call.

What we do with them:

1. We change the setting.
2. We **re-run the full 14,565-JD baseline** and show you the before/after — no decision lands on an estimate.
3. We record your ruling with **who** decided, **when**, and **the reason**. The system will not accept a ratification without all three — that is how an HR ruling becomes part of the rulebook's permanent provenance, and it cannot be recorded anonymously.
4. Every disagreement from your pilot becomes a **permanent test case**, so the tool never makes that mistake again.

A decision is recorded even when you keep a setting unchanged. "HR reviewed this and kept it" is a different, stronger fact than "nobody has looked yet" — and every one of the 194 settings currently sits at the second.

Until you rule, we change nothing.